

RUBBER DUCK THURSDAY · JULY 16, 2026

AI that runs the business

OpenClaw agents for the work between the systems: invoices, follow-ups, field handoffs, lab reviews, and logistics exceptions.

Brad Groux

with Andrea Griffiths

GitHub · 1 PM Central / 2 PM Eastern

RUBBER DUCK THURSDAY

Where AI gets useful after the demo ends

Brad Groux is CEO of Digital Meld and a Microsoft integrations maintainer on OpenClaw. Andrea Griffiths brings the GitHub builder perspective.



The question

What does an agent system look like when it meets a real workflow?



The answer

Context, an operating standard, hard checks, and a person who owns the decision.

Already in practice: OpenClaw carried the business context. Codex delivered the operating surface. [See the Power Apps build.](#)

Rubber Duck Thursday



THE ADOPTION REALITY

PEOPLE

AI success in business is still a people problem.



Trust

If teams do not trust the source, they will not trust the answer.



Reality

The documented workflow and the real workflow are rarely identical.



Ownership

Someone still has to decide, support, and improve the system.

THE MARKET IS TELLING US SOMETHING

6,000

people in Microsoft's new AI implementation unit

Backed by a reported \$2.5 billion commitment to embedded engineering, industry knowledge, change management, and measurable outcomes.

OpenAI

Anthropic

Microsoft

Amazon

Consulting firms



Model access is not adoption.

A license cannot map the workarounds, trust boundaries, or source-of-truth fights.



Implementation is becoming the product.

The valuable work is fitting AI around a company's people, data, and judgment.



The outcome belongs to the business.

The tool is not the religion. The useful, supported workflow is.

START WITH THE WORK SURFACE

Build the harness before you ask for magic.

A serious agent task should carry five boring fields. They move human judgment earlier and make the output reviewable.



01

Mode

Prototype, standard, or production?



02

Context

Which workflow, source files, owners, and systems matter?



03

Boundary

Least privilege per task: what may it read, write, send, or reach?



04

Verification

Which command, check, evidence, or rubric owns correctness?



05

Handoff

What changed, what passed, and what risk remains?

COMMUNICATION IS INFRASTRUCTURE

A center of excellence is a communication network, not a committee.

Bring the people who know the workflow, the data, the risk, and the support burden into the same loop.

Listen



Decide



Document



Review



Improve



Executive sponsor

Sets the outcome, removes blockers, and accepts business risk.



Workflow owners

Know how the work actually moves and where it breaks.



Domain experts

Carry the judgment calls that rarely make it into process diagrams.



Data, IT, security

Define sources, access, controls, observability, and support.



Builders and champions

Prototype, test, teach, collect feedback, and keep the standard alive.

THE SHARED LANGUAGE

Markdown turns context into an artifact everyone can inspect.



Readable by people and agents

No special editor or opaque system required.



Searchable, diffable, reviewable

History, ownership, and rollback come with the repo.



Portable across the workflow

Use the same source for discovery, SOPs, PRDs, prompts, checks, and handoffs.

```
# Accounts Payable Review
```

```
## Trigger
```

```
- Invoice enters the queue
```

```
## Required checks
```

```
- [ ] Purchase order exists
```

```
- [ ] Receipt matches quantity
```

```
- [ ] Duplicate scan passed
```

```
## Stop condition
```

```
Do not approve or pay.
```

```
[Open the SOP](demo/scenarios/accounts-payable/SOP.md)
```

SKILLS ARE OPERATING PROCEDURES

If you repeat an instruction twice, turn it into an SOP.

Chat is a good place to work through an idea. A repo is a better place for the standard.



Purpose and trigger

Why this workflow exists and what starts it.



Steps

The repeatable path, written at the level a new operator can follow.



Verification

The evidence and deterministic checks required before handoff.



Prerequisites

Inputs, access, owners, and source-of-truth records.



Stop conditions

What the agent must never guess, send, approve, or change.




Handoff

Who reviews, who approves, and what is preserved for the next run.

MAKE THE DECISIONS DURABLE

A PRD forces the questions the demo skips.

\$grill-with-docs interviews the plan until the user, owner, source of truth, boundary, proof, and support model are explicit.



BOUND THE AGENT

The model proposes. Deterministic tools judge.



Context

Read the real workflow, artifacts, owners, and constraints.



One bounded action

Draft, classify, patch, or propose the smallest useful move.



Hard verifier

Run calculations, schema checks, tests, policy rules, or source reconciliation.

Invalid



Correct the action

Blocked



Ask for evidence

Verified



Human review

Stop condition: stop after repeated no-gain attempts, missing authority, or missing evidence. Report the best verified result.

REAL-WORLD USE CASE 1

Accounts Payable

The agent assembles the three-way-match evidence and exception queue. A finance reviewer still owns approval and payment.

Fictional data · Draft only · No external actions



Invoice

Read vendor, PO, quantity, amount, and date.



Reconcile

Compare PO, receipt, prior queue, and policy.



Classify

Ready for review, hold, or needs evidence.



Approve

Authorized AP reviewer decides outside the agent.

Demo proof: one clean match, one quantity mismatch, and one likely duplicate. The agent drafts questions and sends nothing.

REAL-WORLD USE CASE 2

Marketing

The agent turns approved facts and reusable assets into a campaign packet with a visible claim ledger.



Brief

Audience, event facts, source assets, and constraints.



Draft

Email, social post, landing-page outline, and open decisions.



Verify

Dates, names, calls to action, and evidence for every claim.



Review

Marketing and legal approve before publication.

Demo proof: an unsupported performance percentage is removed instead of polished into a confident claim.

REAL-WORLD USE CASE 3

Operational agents make constraints visible before work moves.

Field Operations

Draft the morning dispatch handoff without assigning crews or bypassing safety.

Read	Work orders, crew skills, equipment, access windows, permits.
Flag	Permit hold, expiring inspection, late parts, unavailable truck.
Approve	Dispatcher and safety owner make the assignment.

Lab Management

Assemble sample-release evidence without changing or releasing a result.

Read	Sample, method, controls, duplicates, calculations, custody.
Hold	QC failure, method mismatch, incomplete result, custody gap.
Approve	Qualified analyst and lab manager own release.



REAL-WORLD USE CASE 4

A supply-chain agent should run an exception desk, not the company.

Rank risk, show the evidence, and give the owner a bounded next decision. Keep clean shipments quiet.

SHIPMENT	PRIORITY	EVIDENCE	HUMAN DECISION
SH-5204	CRITICAL	Temperature exceeded 25 C for 18 minutes.	Quality owns product disposition.
SH-5207	HIGH	Job-critical material is projected 8.5 hours late.	Operations confirms impact; logistics evaluates approved options.
SH-5201	CLOSED	Delivered within plan at 10:42 AM.	No action. Suppress from the active queue.

The pattern is the product: source packet → SOP → PRD → agent draft → verifier → human approval → feedback.

START MONDAY

Pick one painful workflow. Ship one reviewable artifact.



Map and baseline the real workflow

Find the operators and source of truth. Record cycle time, touch time, exceptions, and rework.



Give the agent a harness

Markdown context, an SOP, a PRD, hard checks, and stop conditions.



Keep the decision human

Start with drafts, review the evidence, and improve the loop with the operators.



github.com/BradGroux



twitter.com/BradGroux



linkedin.com/in/bradgroux



bradgroux.com



github.com/BradGroux/ai-dev-days



Join the SSTB.ai Discord

discord.gg/Gmfkm7QVSF