

Own your operating model. Rent your tools.

The next project should start further ahead.

Brad Groux

Co-Founder · CEO · Builder · Speaker · AI Educator & Evangelist
OpenAI Certified API Expert · ChatGPT Expert · Codex Expert

Purpose Decisions Procedures Evidence Memory

FOLLOW

 @bradgroux

 @bradgroux

 @bradgroux

VISIT

bradgroux.com

digitalmeld.io

sstb.ai

THE PRACTITIONER VIEW

Three years of research. Products in use.



Brad Groux

Digital Meld co-founder
Buzz and OpenClaw Contributor
Veritas Kanban creator

Research

AI-Native Operating Framework

Delivery

Veritas Kanban and OpenClaw
contributions

Teaching

AI Dev Days and Start Small,
Think Big

THE COST OF STARTING OVER

AI can lose the context that makes your work yours.

A new conversation may not know your team's decisions, your definition of "done," or why you rejected an earlier idea.

You repeat the background.

Who the work serves, what matters, and what has already been tried.

The same mistakes return.

A discarded idea comes back. An old assumption becomes a new answer.

The next person starts cold.

The result survives, but the reasoning and unfinished work stay in your chat.

For example: your team decided to exclude pending sales from the weekly revenue report. A new chat includes them again because it never received that decision.

WHAT THE THEME MEANS

Keep your way of working. Change the tools when you need to.

Your operating model is how work gets done: what you are trying to achieve, who decides, how you check quality, and what you learn.

What you should own

Requirements, procedures, decisions, sources, and the record of what happened.

Keep them in formats your people and tools can read.

What you can replace


The model, chat app, coding assistant, or task-management tool.

A better tool should not require rebuilding your organization's knowledge.

For example: an expense report still needs receipts and a manager's approval when you switch accounting software. Your rules belong to your business; the software helps carry them out.


A model is the trained system that generates an answer.

An AI model learns patterns from training data. When you give it input, it uses those learned patterns to generate an output. A language model generates text one piece at a time.



An agent uses AI to take actions toward a goal.


An AI agent uses a model to decide what to do next, uses tools to do it, and responds to the results. It can carry out several steps instead of stopping at an answer.



The difference: a chatbot can explain how to prepare a report. An agent can read the source files, build it, check it, and save the result. It still needs limits and a human owner.

A harness is the software that lets a model do useful work.

The harness connects the model to its environment. It supplies context, makes tools available, controls what actions are allowed, and returns results so the model can continue.



For example: Codex gives a model access to project files and development tools. The model proposes an action; the harness executes it within the configured permissions.

Aggie Launchpad: from an idea to a useful experiment.

Jordan wants to start Campus Cycle: a bicycle-repair pickup service for students. Before building the business, Jordan needs to learn whether students will actually use it.

The app we will build

A founder dashboard connecting assumptions, customer evidence, a pilot experiment, and next decisions.

What we know so far

20 interviews. 12 students expressed interest. Six reserved a free pilot slot. Nobody has paid.

The decision ahead

Run a small pilot if at least five reservations remain active. Jordan decides whether to proceed.

The challenge: two students cancel. Does the app still say “ready to launch,” or help the founder understand what changed?

An attractive startup dashboard can still tell the wrong story.

"Build an app to help me validate my startup idea" is easy to ask. The useful result depends on what the app treats as evidence.

Interest $\neq$ a reservation

A positive interview response does not mean a student booked a pilot slot.

A reservation $\neq$ a sale

A free pilot can test participation. It cannot prove willingness to pay.

A target $\neq$ approval

Meeting five reservations allows a founder review. It does not authorize a launch.

A useful dashboard makes uncertainty visible instead of turning every positive signal into "validated."

WHAT USEFUL OUTPUT SHOULD ANSWER

Can the founder make a better decision?

What do we know?

Show interview interest and active reservations separately, with their sources.

What are we testing?

At least five active reservations for a free pickup pilot. Paid demand remains untested.

What happens next?

Recalculate after cancellations, flag a missed target, and name the decision owner.

The test of the app: can another teammate explain the evidence and next decision without reading Jordan's entire chat history?

A MAP, NOT A LEADERBOARD

Different tools. Different jobs.

Agent work

Codex · Claude Code · Grok Build · DeepSeek Harness

Model execution

Ollama

Personal agent environment

OpenClaw

Collaboration + work state

Buzz · Veritas Kanban

Claude Code

 Read

Understand files and project context.

 Act

Edit code and run permitted commands.

 Check

Inspect results and iterate.

Claude is a model family. Claude Code is the tool around it.

[Claude Code documentation ↗](#)

MODEL RUNTIME

Ollama

Ollama runs and serves models so you can use them from a chat interface, an app, or an agent harness. It supports local models and cloud options.

Why you would use it

Choose where a model runs and connect it to the tools you already use.

[Explore Ollama ↗](#)

Ollama / Aggie Launchpad

Choose a model

Select a local download or a cloud-backed model.

Ask about the source

"How many active pilot reservations do we have?"

Use the response

Your app or harness receives the model's output.

Illustrative workflow

CODING-AGENT HARNESS

Grok Build

Grok Build is a coding agent that works with a project's files and development tools. It can help turn a request into changes you can inspect.

Why you would use it

Build or revise software while keeping the project, instructions, and results together.

[Explore Grok Build ↗](#)

Grok Build / Aggie Launchpad

Understand the project

Read the founder brief and existing app.

Make a change

Add a warning when reservations fall below the target.

Check the result

Verify the reservation count and inspect the changed files.

Illustrative workflow

PERSONAL-AGENT ENVIRONMENT

OpenClaw

OpenClaw connects AI models, tools, and communication channels so an assistant can work with information and perform configured tasks.

Why you would use it

Bring recurring work to an assistant that has the right context and access.

[Explore OpenClaw ↗](#)

OpenClaw / Aggie Launchpad

Ask a useful question

"What do we know about customer demand?"

Consult the context

12 expressed interest; six reserved. Paid demand is untested.

Return a next action

Jordan needs to review the pilot evidence.

Illustrative workflow

Buzz

Buzz is a shared workspace where humans and AI agents collaborate around conversations, projects, and artifacts.

Why you would use it

Keep a team's questions, work, and review in a place where someone else can pick them up.

[Explore Buzz ↗](#)

Buzz / Aggie Launchpad

Jordan · founder

"Here is the updated experiment."

Agent · review

"Four reservations remain. Jordan must review the pilot."

Next person · continuation

The plan, decision, and next owner remain together.

Illustrative workflow

DEVELOPER PREVIEW

DeepSeek Harness

DeepSeek Harness is an extensible agent environment built around plugins. Models, tools, sessions, and other capabilities can be composed through configuration.

Why you would use it

Explore how an agent is assembled and inspect the record of what it actually did.

[Explore DeepSeek Harness ↗](#)

DeepSeek Harness / Aggie Launchpad

Compose the agent

Choose the model, tools, and operating mode.

Carry out the work

Read the brief and check the experiment evidence.

Inspect the trajectory

See the context, tool calls, and results.

Illustrative workflow

A SHARED VIEW OF THE WORK

Veritas Kanban

Veritas Kanban is a local-first task board with optional AI agent orchestration. It makes ownership, status, and completion criteria visible.

Why you would use it

Coordinate human and agent work without relying on a chat message to tell you what is done.

[Explore Veritas Kanban ↗](#)

Veritas Kanban / Aggie Launchpad

To do · Jordan

Review the pilot reservation evidence.

Blocked · Jordan

Decide what to do after two cancellations.

Done means evidence

Attach the source record and accepted decision.

Illustrative workflow

WHAT THE NEXT CONVERSATION NEEDS

Memory keeps the reasoning with the work.

 Current facts

Six active reservations at the start. They are for a free pilot, not paying customers.

 Decisions and reasons

Jordan chose five reservations as the minimum to review a pilot. This is a local experiment target.

 Open questions

Will students pay? Can pickup and repairs be delivered reliably? Neither has been tested.

A useful memory: "Do not describe interview interest as sales. Preserve the source counts and the reason for each decision."

What are we building?

How do we do it?

PROJECT REQUIREMENTS DOCUMENT

Define the result.

The dashboard must distinguish interest, reservations, cancellations, and paid demand.

Show active reservations and compare them with the five-reservation target.

+

STANDARD OPERATING PROCEDURE

Define the work.

Read evidence, apply updates, check counts, identify exceptions, and save the handoff.

Target missed? Flag it and ask Jordan to decide. Never invent approval.

A WAY TO DEFINE HOW PEOPLE AND AI WORK TOGETHER

The AI-Native Operating Framework



Intent

What outcome are we responsible for?



Responsibility

Who owns the work and decisions?



Work

What inputs, actions, and handoffs are needed?



Control

What limits and exceptions apply?



Assurance

What evidence makes the result acceptable?



Learning

What do we retain and improve?

Read the [open framework](https://github.com/BradGroux/ai-native-operating-framework) → github.com/BradGroux/ai-native-operating-framework

Example: two cancellations leave four active reservations against a target of five. The agent flags the gap; Jordan decides whether to recruit, revise, or pause. Save the decision and its evidence.

PUTTING THE FRAMEWORK INTO PRACTICE

Give the next experiment something useful to inherit.

 Requirements

Keep interview interest, reservations, and paid demand separate. Show sources and uncertainty.

 Procedure

Recount when evidence changes. Compare with the target. Route the next decision to the founder.

 Memory

Retain the hypothesis, target, source records, decision history, and unanswered questions.

A reusable rule: “Never turn an expression of interest into a paying customer.” That correction should survive the next project and the next tool.

Make the evidence visible.

Keep the assumptions visible, too.

Observed in the scenario

20 interviews, 12 expressions of interest, six initial free-pilot reservations.

Still an assumption

Students will pay for pickup. The team can deliver repairs at a sustainable cost.

A bounded experiment

Use five active reservations as the threshold for reviewing a small free pilot.

What changes in the app: separate evidence cards, an explicit experiment target, and a pending founder decision. No blanket “business validated” badge.

THE EXPERIMENT CHANGES

New evidence should change the next decision.

Two students cancel

Six initial reservations – two cancellations



Four remain active

Below the target of five



Jordan reviews

Recruit more, revise the experiment, or
pause

What the app should do: update the count and status, retain the original six, and flag the decision. Twelve interested interviewees do not fill the missing reservation.

A HANDOFF THAT SUPPORTS A DECISION

The next teammate needs the learning, not just the app.

"Students love the idea."

The next teammate cannot tell who committed, what changed, or whether anyone tested payment.

"Four reservations remain."

Two canceled. The target is five. Jordan must decide what comes next. Paid demand is still untested.

Save with the handoff: the hypothesis, source records, original and current counts, target, decision owner, and next unanswered question.

Evidence for the next decision.

Aggie Launchpad / Campus Cycle

Experiment: free bicycle-repair pickup pilot

Interview interest

12 / 20

Interest reported · not sales

Active reservations

4 / 5 target

Six initial – two canceled

Decision needed

Target missed

Jordan · review pending

Evidence and assumptions

Observed Four active reservations for a free pilot.

Untested Whether students will pay for pickup and repairs.

Next choice Jordan: recruit, revise, or pause. No decision recorded.

Handoff for the next teammate

Two reservations were canceled. The five-reservation threshold has not changed.

Do not describe the venture as validated. Review the pilot before proceeding.

Saved with the experiment

Sources · counts · target · decision history

WHY THIS MATTERS BEYOND ONE DEMO

The value compounds when the knowledge stays.

First project

Clarify the rules and check the result



Next project

Reuse the rules and accepted lessons



New tool or teammate

Bring the operating knowledge with you

Less re-explaining

Requirements and decisions no longer depend on one chat.

More consistent work

The same approval rules and quality checks travel with the task.

Easier change

Replace a tool without discarding what your team has learned.

A PRACTICAL FIRST STEP

Start with one decision you want AI to help you make.

Choose a customer experiment, research summary, weekly report, or project handoff. Write a short working agreement before asking AI to help.

What counts as evidence?

"Separate customer behavior from what people say they might do."

Who decides?

"The founder accepts the next experiment. AI proposes and checks."

What should survive?

"Save the sources, decision, reasoning, and next question."

Your first move: write those three answers for your own task, try them with your current tool, and retain the corrections.

FOR NEW MEMBERS ATTENDING TODAY

Your first year of SSTB.ai is on us.



**Start Small,
Think Big.**


Keep learning, building, and sharing
with the SSTB community.

No payment card required.

365 days of membership access.

Redeem by **September 16, 8 PM Central.**

Use the email approved at event check-in.



Scan to claim your year.

New members · approved attendees only

YOUR NEXT MOVE

Make the next project start further ahead.

-
- | | |
|---|------------------------------|
|  Pick one workflow. | Choose something you repeat. |
|---|------------------------------|
-
- | | |
|--|--|
|  Name the owner. | Make decisions and authority explicit. |
|--|--|
-
- | | |
|--|--|
|  Define one check. | Say what evidence will make you accept the result. |
|--|--|
-
- | | |
|---|---|
|  Keep the lesson. | Save the context the next person needs. |
|---|---|
-

RESOURCES AND FOLLOW-UP

Keep building.

AI Dev Days

Presentation and learning materials

AI-Native Operating Framework

The open operating framework

SSTB.ai

Start Small, Think Big

Brad Groux

Connect on LinkedIn

Digital Meld

Research and product delivery



Start with the source.

github.com/BradGroux/ai-dev-days

FOLLOW

 @bradgroux

 @bradgroux

 @bradgroux

VISIT

bradgroux.com

digitalmeld.io

sstb.ai